

OpenLCB Technical Note

Event Transport

Oct 19, 2023

Draft

1 Introduction

This Technical Note contains informative discussion and background for the corresponding OpenLCB Event Transport Standard. This explanation is not normative in any way.

OpenLCB nodes respond to local inputs and state changes by emitting Producer/Consumer Event Report (PCER) messages. When they do this, they are acting as “Producers”. OpenLCB nodes receive these messages, and can act on them locally if desired. If they do so, they are acting as “Consumers”. Together, this is called the Producer-Consumer Model of system control. In this model a particular concept, or event, is the sum of the inputs that trigger the production of the PCER by the producer(s) and the action(s) taken by the consumer(s).

Its main advantage is the knowledge independence it imparts between all of the nodes involved in an event. Expressed in more anthropomorphic terms, neither the producers nor consumers of a PCER are aware of each other. This adds flexibility to the system, since more producers or consumers of the PCER can be added to implement the associated event, without concern of side effects to the other nodes involved in the event.^{1,2}

The Producer/Consumer Event ID included in a PCER message is an 8-byte unique identifier. The method for assigning these and ensuring their uniqueness is in the OpenLCB Event Identifiers Standard and Technical Note.

There are many possible ways to use these 64 bits, some of which are discussed as examples below. In particular, OpenLCB does not require or enforce any particular partitioning of the eight bytes. More than one node may emit the same Producer/Consumer Event ID (note that the Node ID of the emitting node is available separately in the message). More than one node may receive and act on a PCER message with a specific Producer/Consumer Event ID. The only requirement is that the 8-byte quantity be unique to the event.

In addition to the 8-byte EventID, produced events may be accompanied by one or more bytes of payload. This can be used as a global transport mechanism.

OpenLCB event IDs carry no specific information about whether they are coding “on” or “off”, “active” or “inactive” for some output state. It's possible for multiple consumers to do different things in response to a single event: A event that means “set nighttime lighting” might turn some outputs off and turn others on. Simple nodes will often use two producers to represent “input line N went on” and “input line N went off”; they will often use two consumers to represent “turn on output line N” and “turn off output line N”.

¹D. Miorandi and S. Vitturi, “Performance analysis of producer/consumer protocols over IEEE 802.11 wireless links”, *Proc. IEEE Workshop Factory Communication Systems (WFCS)*, 2004

²The Producer/Consumer Model and Control System Design ControlLogix 1999

https://cours.etsmtl.ca/gpa774/Cours/old-24-03-04/Documentations/Rockwell/articles/Producer_Consumer.html

OpenLCB events therefore do not directly represent states. Rather, they represent state transitions, or rather a change to the state of the layout. One event might be interpreted as “go to output N on”, another as “go to output N off”, and only by seeing which transition happened most recently can the state of output N be recovered.

2 Annotations to the Standard

This section provides background information on corresponding sections of the Standard document. It is expected that these two documents will be read together.

2.1 Introduction

Nothing to add to the Standard.

2.2 Intended Use

A particular PCER may be produced in response to physical input or state change on a layout, for example a contact closure or the activation of a block occupancy detector, or it can be in response to a non-physical change such as the arrival of a certain time. The same PCER can be produced by several different state changes. That PCER may result in some action at consumers. The term 'event' includes the state changes, the PCER message, and the resulting actions, and represent a specific idea or concept, such as “Night has fallen”.

For example, the concept of an “Emergency stop” event might include the actions of turning off track power, turning on a fault-indicator, and sounding a buzzer. The PCER may be produced by any of: a panic button, a throttle button, an over-current detector, or an anti-collision system.

The aforementioned “Night has fallen” event might be triggered by a fast-clock time, a toggle switch, or a control program, and its actions might include illuminating street lights, house lights, and the dimming of overhead lighting.

In addition, ranges of PCERs can be used for specific purposes. For example, a fast clock's messages can be implemented as a range of PCERs, where the low order bytes code the time.

2.3 References and Context

For more information on format and presentation, see:

- OpenLCB Common Information Technical Note

For more information on MTI values and how they are laid out, see:





- OpenLCB Message Network Standard
- OpenLCB MTI allocation table³.

2.4 Message Formats

The range mask is implemented as a combined base value and mask in order to save space. The base value takes the form 0bv...0 and the mask 0b0...m. If the lowest bit of the base value is a one, then the

³<http://openlcb.org/trunk/specs/MtiAllocations.pdf>

- 65 combined form is $0bv\dots 0$. However, if the lowest bit of the base value is a zero, then the mask is inverted and the combined form becomes $0bv\dots 1$.



- 70 Ranges can only be specified in powers of 2, i.e., they can specify ranges of 2, 4, 8, Therefore, if a range of n is required, then the next larger power of two is used. Ranges should not be over specified, so the standard requires that a range not extend past the next power of two by requiring that at least 50% of the EventIDs in the range be relevant.

- 75 Encoding a range is done by determining how many values are required, choosing the next higher power of two, choosing a base value and the mask, and finally combining the two, inverting the mask if necessary. For example, a range $0x1234xx$ would require a mask of the form $0x0000FF$, and the combined form would be $0x1234FF$. A range of $0x1235xx$ would also require a mask of $0x0000FF$, but the combined form would be $0x123500$, since the mask is inverted. In this example, an event range of 256 events and a 16-bit address is specified.

Decoding a range involves identifying the mask as lowest contiguous bits of the same value as the lowest bit. The base value is the remaining high order bits. For example, the combined form

- 80 0x987650 would yield a mask of 0x00000F (inverted) and a base of 0x98765, whereas 0x98764F would yield a mask of 0x00000F and a base of 0x987640. In this example, an event range of 16 events and a 20-bit address is specified.

2.4.1 Producer/Consumer Event Report (PCER)

- 85 This message transports an Event ID from a producer node(s) to zero or more unspecified consumer nodes. It is the backbone of the protocol, and forms most of the expected traffic.

The source Node ID allows identification of which node sent the event. This can be used for diagnostic and status purposes, but nodes must act on the event without regard to where it came from.

- 90 Note that the Standard limits the payload to 256 bytes. This is to put a bound on required buffers. In practice, very long PCER-with-payload messages can block timely transmission of e.g. Datagram and Stream messages and other PCER-with-payload messages. When designing protocols that use this capability, it may be useful to define ways to use multiple, shorter PCER-with-payload messages instead of one long one.

2.4.2 Identify Consumer

This message is broadcast and requests each node to report whether it consumes this Event ID.

- 95 **2.4.3 Consumer Identified**

This message is broadcast, in response to a received Identify Consumer message, from each node that consumes the included Event ID. Nodes also send this as part of their startup so that other nodes know they want to consume particular events.

This is one of the messages that allows bridges to do automatic routing of event messages.

- 100 The “valid”, “invalid”, “unknown” subforms are used to attempt to recover the state of the overall system. In general, event transmission via PCER messages conveys state changes. If a node comes in during operation, it might want to determine the correct value of the distributed state. Producers can, but don't always, know that state. For example, consider “On” and “Off” events for a light. There might be two producers attached to a toggle switch. In that case, one of them will still see their state as current, and will reply with valid. Alternately, the two producers might be set to produce when two separate pushbuttons are pressed. In that case, at some later time neither producer knows whether the distributed state corresponds to its event. They both have to report as unknown. Consumers can also know the state. The lamp controller that consumes the “On” and “Off” probably knows whether it is currently providing power to the lamp, and so can reply with valid for one event and invalid for the other.

2.4.4 Consumer Range Identified

This message broadcasts, in response to a received Identify Consumer message, from each node that consumes events in the range specified by the included event-ID-with-mask. This is one of the messages that allows bridges to do automatic routing of event messages.

- 115 The requirement that 50% or more of the Event IDs in a range are consumed is to allow rounding the range up to the enclosing bit, but no further. This keeps the Event ID ranges in the complete OpenLCB system compact and is intended to make the task of gateways more manageable.

Allowing this requirement to be overridden by other Standards permits development of standard protocols that use one Consumer Range Identified message for separate disconnected ranges to

- 120 improve efficiency. By only allowing this to be done via the Standards process, not as an ad-hoc design issue, complexity is kept under control.

2.4.5 Identify Producer

This message is broadcast and requests each node to report whether it produces this Event ID.

2.4.6 Producer Identified

- 125 This message is broadcast, in response to a received Identify Producer message, from each node that produces the included Event ID. This is one of the messages that allows bridges to do automatic routing of event messages.

The ProducerIdentified message is different from a Producer-ConsumerEventReport message for several reasons:

- 130
- It allows the valid/invalid/unknown subforms, which don't make any sense for a PCER message
 - It allows network listeners to separate the event identification mechanism, which uses ProducerIdentified, from the event transport mechanism, which uses PCER. Otherwise, nodes might think that an event had occurred and was being produced, when it was just a response to a request for current status.

135 2.4.7 Producer Range Identified

The requirement that 50% or more of the Event IDs in a range are produced is to allow rounding the range up to the enclosing bit, but no further. This keeps the Event ID ranges in the complete OpenLCB system compact and is intended to make the task of gateways more manageable.

- 140 Allowing this requirement to be overridden by other Standards permits development of standard protocols that use one Producer Range Identified message for separate disconnected ranges to improve efficiency. By only allowing this to be done via the Standards process, not as an ad-hoc design issue, complexity is kept under control.

- 145 An example of a producer range would be a fast clock node that sends a range of events to indicate the time. Since there are $24 \times 60 \times 60 = 86,400$ seconds in a day, it needs a range containing 86,400 events. In this case, rather than having to send 86,400 individual ProducerIdentified messages at start-up or on request, it can send one Producer Range Identified message with the appropriate base/mask Event ID Range value.

- 150 This message is broadcast, in response to a received Identify Producer message, from each node that produces events in the range specified by the included event-ID-with-mask. This is one of the messages that allows bridges to do automatic routing of event messages.

2.4.8 Identify Events

Identify Events comes in both addressed and global forms. They are intended to be used for different use cases.

- 155 In some cases, the node may want to query a specific node as to what events it produces or consumes. In these cases the directed form is appropriate. This message will minimize the load on the network as only the node of interest will respond.

In other cases, the node producing the Identify Events message has a lack of knowledge to which nodes produce or consume the events of interest, and therefore the global form is the more appropriate choice.

160 The global form of Identify Events is similar to the combination of an Identify Producers and an
 165 Identify Consumers message for every possible enumeration of Event Ids. This event is useful to build
 routing tables for networks segments or network browsing tool but be aware that since this form is
 transmitted throughout the network, it can be relatively expensive in terms of bandwidth.

2.4.9 Learn Event

165 This message is part of a communication-based method for configuring various producers and
 consumers across the OpenLCB network to respond to a single Event ID. After those producers and
 consumers have been selected via some other process, this message is broadcast to carry a specific
 Event ID to all of them.

2.5 States

170 After the Initialization Complete message is sent, but before any -Producer/Consumer Event Report
 messages are sent, each node must identify all events produced or consumed via zero or more Identify
 Consumer, Identify Consumer Range, Identify Producer and Identify Producer Range messages. These
 are not required to be in any particular order.

2.6 Interactions

175 Nothing prevents extra Producer/Consumer Identified messages. Nothing prevents combining replies to
 multiple requests. This allows simplified implementations, for example setting a bit to indicate that a
 reply can be sent when time/priority is available.

Nodes can send identify/identified messages for automatically-routed Event IDs. Nodes aren't required
 to reply to Identify Producer/Consumer messages for automatically-routed Event IDs to simplify node
 programming, but it is useful if they do.

180 Identify Producer/Consumer messages for well-known Event IDs require response; there's no
 exemption for those.

The state of the system can sometimes be constructed from the received Producer/Consumer Identified
 messages and their "valid", "invalid", "unknown" subforms. It is possible that these will result in
 conflicting and/or ambiguous state.

185 Note that the Advertised/Not-Advertised state machine is reset by -the Initialization Complete message,
 and not by any lower level link (CAN or other) state machine.

Delay in sending the Producer Identified / Producer Range Identified and Consumer Identified /
 Consumer Range Identified messages after Initialization Complete is allowed, but there's no delivery
 guarantee for this node's events until those have been sent and the state has transitioned to Advertised.

190 When a node changes the events it manages, it still needs to emit the Producer Identified / Producer
 Range Identified and/or Consumer Identified / Consumer Range Identified messages. For example,
 reconfiguration could cause this. It can be handled by sending individual messages, or by resetting and
 sending them all as part of that process.

195 There is no way to indicate that a node is no longer interested in a particular Event ID. (Sending
 Initialization Complete again says that all events are not interesting, though). This could be added, but
 it is much harder for gateways to decide that an Event ID is not interesting than that it is interesting,

because they have to keep a list of all the node IDs that are interested, and back that off. Better to just let the set of interesting Event IDs grow until the layout or gateway is reset.

200 To ensure that event messages are properly routed, nodes must announce when they start to produce or consume messages. Specifically, they must do this at two times:

- When the node is first initialized, after the “Initialization Complete” message has been sent.
- When a configuration change has added another Event ID that can be produced or consumed.

205 To announce new Event IDs, the node must transmit a Producer Identified message for each Producer/Consumer Event ID it can newly produce, and a Consumer Identified message for each Producer/Consumer Event ID that it can newly consume.

2.6.1 Event Transfer

The last sentence is the key to gateway traffic reduction. See §3.1 below.

2.6.2 Event Inquiry

210 Identify Events comes in both global and addressed forms. The global form can place a very large load on the network and should only be used when needed.

The Identify Events message is sent to request that the specified nodes report all the events they produce or consume. These reports can be Identified Event messages specifying individual Event IDs or ranges of events.

215 Identify Events is useful to be able to rapidly determine which, if any, Producer/Consumer Event IDs that a particular node is listening for and that it can emit. This can be used as a configuration diagnostic.

To determine which Producer/Consumer Event IDs can be sent by a particular node, the inquiring node sends an Identify Events message addressed to the target node.

220 The node must reply with a Producer Identified message for each Producer/Consumer Event ID it can produce, and a Consumer Identified message for each Producer/Consumer Event ID for which it is listening.

2.6.3 Producer Inquiry

225 It is useful to be able to determine which, if any, nodes are configured to possibly transmit a specific PCER message and the current state of those producers. This can be used as a configuration diagnostic, and as a way of building filtering and routing tables.

To determine which nodes can send a particular Producer/Consumer Event ID, a node sends an Identify Producers message carrying the desired Producer/Consumer Event ID. This is an unaddressed message addressed to all nodes.

230 All nodes that are listening for that Producer/Consumer Event ID reply with a Producer Identified broadcast message.

The “valid” bit indicates that the node's internal condition is consistent with sending this Producer/Consumer Event ID. For example, assume a node sends Producer/Consumer Event ID 2 when the input goes active and Producer/Consumer Event ID 4 when the input goes inactive. Then if the

235 input was active and the node was asked about Producer/Consumer Event ID 2, it would reply “valid”;
 if asked about Producer/Consumer Event ID 4, it would reply “not valid”. Depending on the node's
 structure, it might not always be possible to set the “valid” bit with certainty, in which case the
 “unknown” bit must be set.

240 You can query a state when you come up with the Identify Events message. Conflicting states can
 happen, and have to be addressed. You might also want to enquire of the consumers if you don't get a
 definitive answer from the producers.

In regard to valid/invalid/unknown, it is always possible to send only unknown, or to ignore the three
 sub-forms, but a lot of capability is then lost.

2.6.4 Consumer Inquiry

245 This message is useful to determine which, if any, nodes are listening for a specific Producer/Consumer
 Event ID message.

This can be used as a configuration diagnostic, and as a way of building routing tables.

To determine which nodes are listening for a particular Producer/Consumer Event ID, a node sends an
 Identify Consumers message carrying the desired Producer/Consumer Event ID. This is an unaddressed
 message processed by all nodes.

250 All nodes that are listening for that Producer/Consumer Event ID reply with a Consumer Identified
 broadcast message.

255 The “valid” bit indicates that the node is currently in the state it would be if this message had been
 received last. For example, assume a node sets its output active for Producer/Consumer Event ID 2, and
 inactive for a Producer/Consumer Event ID 4. Then if the output was active and the node was asked
 about Producer/Consumer Event ID 2, it would reply "valid"; if asked about Producer/Consumer Event
 ID 4, it would reply "not valid". Depending on the node's structure, it might not always be possible to
 set the “valid” bit with certainty, in which case the “unknown” bit must be set.

2.6.5 Teach/Learn Configuration

260 Entry-level OpenLCB users should be able to buy two OpenLCB nodes and configure them to do
 something useful with only the pushbutton(s) and LED(s) on them, without requiring a separate
 configuration tool, computer, etc.

265 Producers (in nodes that connect to buttons, toggle switches, occupancy detectors, etc) and consumers
 (in nodes connected to signals, LEDs, displays, layout lighting, etc) need to be configured to use a
 common EventID, such that activating a producer will send the common EventID, which will cause
 one or more consumers to activate.

The Teach/Learn Configuration interaction provides communications support for that. It works without
 any central configuration mechanism, nor even any global knowledge of which producers and
 consumers are being configured.

270 The configuration starts by using external mechanisms, e.g. walking around and pushing buttons, to
 select producers and consumers which will learn the common EventID. These are colloquially called
 "Learners". The Standard does not discuss how a node defines its local user interface for doing that

selection process⁴. A single producer or consumer then provides the EventID that is to be learned in a “Teach Event” message. If the taught Event ID is a new unique ID, then only the selected producers and consumers will share it. If producers and/or consumers are already using that Event ID, the newly configured producers and/or consumers will be tied into the existing network of producers and consumers.

2.6.5.1 Blue/Gold Buttons Implementation Example

General Description

Nodes can be built with two buttons, called Blue and Gold, which can be used to configure the nodes' hardware, including event IDs from one node to one or more other nodes.

- On one or more 'learner'-node(s), using the Blue-button, choose a consumer/producer, and then push the Gold-button to set it to learn. This can be repeated on the same node or on others.
- On the 'teacher'-node, push the Gold-button to choose teach-mode, then using the Blue-button, choose a consumer/producer.
- Push the Gold-button to send the Learn Event message.
- As a result, all producers selected above will activate all the consumers.

Summary:

- Select to learn: blue, blue, ..., gold.
- Select to teach: gold, blue, blue, ..., gold.

Notes:

- Pushing the Blue-button selects consumer/producer's in a circular fashion, with a 'no-selection' before starting again at the beginning
- The teaching- producer/consumer can also be a 'learner'.
- New teaching will replace the previous teaching.
- Pushing Gold for more than 3-seconds will reset the node to its factory events.

Layout Example

Given two Button-nodes, each with three buttons, A-C and D-F, respectively, and four Signals-nodes, 1-4:

1. Buttons set Aspects: Push blue on Signal-1 until the 'green' aspect is selected, then push gold. On the first Button-node, push gold, and then blue once to select button-A, then push the gold again. On Signal-1 use the blue-button to select aspect 'red', and push gold, and then on the Button-node push gold and the blue twice to select button-B, and then gold again. Now, pushing button-A will change Signal-1 to green, and button-B will change it to red.
2. Make a duplicate button: On the second Button-node, push blue once to select button-D, and then gold, and on the first Button-node, push gold, then blue once and then gold. On the second Button-node, push blue twice to select button-E, and then gold, and on the first Button-node,

⁴Having conventions for the user interface would make it easier for users to work with multiple OpenLCB products. Some work has gone into a system for doing this with two buttons and two LEDs, called the “Blue/Gold configuration”, see the next section, but this has not become a consensus solution as of the time of this writing.

push gold, then blue twice and then gold. Now, button-D will turn Signal-1 to green, and button-E will turn it to red.

3. Add a second consumer-node: On Signal-2 push blue to select red and then gold, then on Signal-1 push gold, then blue to select green, and then gold. On Signal-2 push blue to select green and then gold, then on Signal-1 push gold, then blue to select red, and then gold. Now pushing button-A sets Signal-1 to green and signal-2 to red, and button-B turns Signal-1 to red and Signal-2 to green.

2.6.6 Resetting to Default

Nodes are responsible for ensuring that an Event ID is really unique if required. This has implications for how event IDs are defined when the node is reset to “factory defaults”.

For example, say that a node has Node ID 1.2.3.4.5.6 and therefore starts with its first producer containing the 1.2.3.4.5.6.0.1 Event ID. As it comes from the factory, this Event ID is guaranteed to be unique because it comes from the node's unique ID.

This Event ID can be then be configured into several other nodes, which are using it for some particular purpose, say “start of day”. The node 1.2.3.4.5.6 producer can then be reconfigured to contain some other event ID, say 6.5.4.3.2.1.99.99 which is used to mean “end of day”.

If the node is reset to defaults and returns to 1.2.3.4.5.6.0.1, that's no longer a “unique” Event ID. It's being used by another part of the system for “start of day”. Node 1.2.3.4.5.6 will not be behaving as expected for a reset node, as it'll be responding to events taking place on the network. Further, it's 1.2.3.4.5.6.0.1 Event ID is not safe to be copied out into other producers and consumers.

One way to do this is to maintain a counter of created unique Event IDs, and use the next set when doing a reset to defaults. There are 65,538 different Event IDs possible for a node, so even one that uses 512 Event IDs can be reset to defaults a lot of times before its are risk of running out.

Alternately, the node can keep track of which Event IDs might have left the node and been reused. This includes any of the messages that carry the Event ID, including but not limited to Teach Event, plus configuration memory reads and perhaps other methods. If the Event ID in a producer or consumer has been retrieved and used, then it needs to be given a new Event ID. If not, it can keep it's existing Event ID. While Identify Event messages can be used to discover if a particular Event is being used, it will not find any nodes that are not active, or have left the layout, so a conflict could occur if those nodes become active of the layout again. It is therefore much safer to reset to defaults and obtain a new set of unused Events. This is particularly important if the node is passed onto someone else.

2.7 CAN Adaptation

For reference, we describe the messages defined in this standard as formatted for the CAN transport protocol. This section is not normative; should there be a conflict between this document and the Standards, then the Standards take precedence.

The CAN Frame Transport Standard and Message Network Standard define how OpenLCB messages are carried across CAN networks. Following those standards, the Event Transport messages used on CAN are as defined in the following table.

MessageName	Simple Node	Dest ID	Event ID	Header Format	Data-part Content
Producer/Consumer Event Report (PCER)	Y	N	Y	0x195B,4sss ⁵	Event ID
<u>Producer/Consumer Event Report With Payload first frame</u>	<u>Y</u>	<u>N</u>	<u>Y</u>	<u>0x19F1,6sss</u>	<u>Event ID</u>
<u>Producer/Consumer Event Report With Payload middle frame</u>	<u>Y</u>	<u>N</u>	<u>Y</u>	<u>0x19F1,5sss</u>	<u>8 bytes of payload</u>
<u>Producer/Consumer Event Report With Payload last frame</u>	<u>Y</u>	<u>N</u>	<u>Y</u>	<u>0x19F1,4sss</u>	<u>Last 1-8 bytes of payload</u>
Identify Consumer	Y	N	Y	0x198F,4sss	Event ID
Consumer Identified	N	N	Y	0x194C,4sss — Valid 0x194C,5sss — Invalid 0x194C,7sss — Unknown	Event ID
Consumer Range Identified	N	N	Y	0x194A,4sss	Event ID Range
Identify Producer	Y	N	Y	0x1991,4sss	Event ID
Producer Identified	N	N	Y	0x1954,4sss — Valid 0x1954,5sss — Invalid 0x1954,7sss — Unknown	Event ID
Producer Identified Range	N	N	Y	0x1952,4sss	Event ID Range
Identify Events	Y	N	N	0x1997,0sss	
	N	Y	N	0x1996,8sss	fddd ⁶
Teach Event	Y	N	Y	0x1959,4sss	Event ID

Note that a PCER message may carry its Event ID in two types of CAN frames with header format 0x19F1,6sss and 0x195B,4sss. Even if a node has no use for the payload of a PCER-with-payload message, it should still examine and react to the Event ID carried with both header forms.

The 0x19F1,6sss, 0x19F1,5sss and 0x19F1,4ss sequence for first, middle and last frame of a PCER with payload message were chosen so that the first frame has lower CAN priority than the middle frame, which has lower CAN priority than the last frame. This ensures that once a node starts transmitting a PCER with payload message, no other node's first PCER with payload frame will be able to arbitrate and appear on the CAN bus while the first message is being transmitted.

The requirement that the frames of a PCER message with payload be sent without inter-frame gaps is meant to remove the possibility that parts of two messages will overlap due to a collision. Imperfect implementations may have gaps and allow such collisions to take place. Implementors should assume

⁵sss — The 12-bit source alias field.

⁶fddd — First two bytes of the data-part, representing the flag bits and 12-bit destination Alias. See the CAN Frame Transport Standard.

that EWP messages may overlap. Although the most probable collision case is that an EWP message will only overlap with one other EWP message, this cannot be assumed in general. Best practice is to allow for variable buffering up to the available node memory.

2.7.1 CAN States

There are no CAN-specific modifications to the states described in Section 5 above.

2.7.2 CAN Interactions

There are no CAN-specific modifications to the interactions described in Section 6 above.

3 Background information

This section is general information of interest to the reader, implementers, etc.

An OpenLCB Event represents the sum of all the causes and actions associated with particular concept, such as “Night is falling”, and implements a many-to-many relationship. It has been referred to as the Producer-Consumer Model, in opposition to the Master-Slave Model.⁷ Each Event is abstracted into a single unique event-number. An Event-message contains only this event-number. An advantage of this model is that new triggers or new causes which cause the production of the event-number, or new actions resulting from the consumption the event-number, can be implemented by a node without reference to other producers or consumers, which makes this model very flexible and extensible.

Events are most often used to advertise changes in system state, for example an occupancy detector's state. The model maintains the distributed-state model of a layout because Events-messages are broadcast⁸ to all nodes, and this allows individual producers and consumers to use them to change their knowledge of the distributed state.

Other important, and perhaps not obvious, uses include: publishing a capability using an event from a list of well-known Event IDs; transmitting well-known ranges of data, such as fast-clock time; and embedding an other system's message into the bytes of the Event ID, for example a DCC accessory message. Notice that these latter uses allows any node to respond to these messages using its usual event-handling. For example, a node can activate an output in response to a specific fast-clock time.

3.1 Gateways

Gateways can route PCERs only to segments with nodes expressing interest by processing the other event messages, including Consumer and Producer Identified.

The automatically-routed event range always has to be routed. This can be done by making a permanent entry in routing tables, or any other implementation method.

Note that gateways may, but are not required to, reorder lower priority messages after higher priority messages. This only applies to OpenLCB priorities, of which there are four values. Gateways are not allowed to reorder messages that have the same priority. This means that they may not reorder 0x5B4's

⁷ For a general discussion, see: <http://openlcb.org/trunk/documents/notes/ProducerConsumerModel.html> and http://openlcb.org/trunk/scratchpads/JohnSL/Tech_Note_Introduction_to_the_PC_Model_2007-10-10.pdf

⁸ OpenLCB uses interest-based routing, which means that Event messages are actually broadcast only to those bus segments on which a node has expressed interest. Event-ranges are used to express interest in a classes of events.

between each other, but they may reorder a 0x5B4 to come before any event with payload message (0xF14). It is not guaranteed that events and events-with-payload messages arrive in the same order as they were sent by a given node. Applications must be aware of this when they implement distributed state machines or controlling logic.

3.2 Examples

Consider a OpenLCB installation with two CAN segments A and B connected via gateways and a TCP/IP link. In addition, a program is attached via a TCP/IP link and additional gateway to CAN segment B.




Figure 1 – Example OpenLCB Network

The program has to be able to fully communicate with all OpenLCB nodes on both segments A and B. These examples show how the gateways can make that happen. For simplicity, the labels on the figure will be used as if they were the full NID of the node, e.g. "GB1" for the left gateway on segment B.

Frames and messages are represented by ["source address", "content", "optional destination address"] triplets. Not all messages are shown. Note that there are many ways that gateways can manage traffic, and this is just one approach.

Before the cases below, both segments come up, assign the numbers 1,2,3 as aliases to their local nodes, 10 to gateway 1 and 20 to gateway 2 on segment B. For simplicity, A1 produces event X, which A2 and B2 want to consume. A1 also produces Y, which no other node wants to consume.

3.2.1 Case 1: All nodes initialized before link established, event X is sent

The A-B gateways initialize their internal tables.

A1 sends X. GA receives it. Because GA hasn't seen X before, X is forwarded across the link to GB, followed by an "Identify Consumers of X" message.

GB puts the X message on the CAN segment, and B2 consumes it.

GB puts the "Identify Consumers of X" message on the CAN segment.

B2 responds with "Consumer of X Identified"

GB receives "Consumer of X Identified", forwards to GA, which marks its tables to continue forwarding these. GA forwards that message on the CAN segment in case there are other listeners for it.

3.2.2 Case 2: Node B2 comes up, after event X has been sent in the example above

The gateways are up, but don't know about B2. GA has seen X sent by A1, but has marked it as "do not forward" because there was no consumer for it on B.

B2 comes up, gets a NID alias assigned, sends “AMD” and “Initialization Complete”.

425 B2 sends “Consumer of X Identified”. GB forwards that to GA.

GA remembers that event X must be forwarded across this link if seen.

A1 sends X, GA receives it, forwards it across link.

3.2.3 Case 3: Node A1 comes up, then event X is sent

The gateways are up, but don't know about A1 or X.

430 A1 comes up, gets a NID alias assigned, sends “AMD” and “Initialization Complete”.

A1 sends “Producer of X Identified”. GA forwards that to GB.

GA notes that event X currently must be forwarded across this link if/when seen.

A1 sends X. GA receives it. Because GA hasn't seen X before, X is forwarded across the link to GB, followed by an “Identify Consumers of X” message.

435 The rest of this proceeds like case 1.

3.2.4 Case 4: Node A1 comes up, then event Y is sent

The gateways are up, but don't know about A1 or Y.

A1 comes up, gets a NID alias assigned, sends “AMD” and “Initialization Complete”.

A1 sends “Producer of Y Identified”. GA forwards that to GB.

440 GA notes that event Y currently must be forwarded across this link if seen.

A1 sends X. GA receives it. Because GA hasn't seen Y before, Y is forwarded across the link to GB, followed by an “Identify Consumers of Y” message.

No “Consumer of Y Identified” message is received back, so after a short timeout GA marks it's tables to not forward Y. (If a not later comes up that consumes Y, a sequence similar to Case 2 occurs.)

445 3.3 Node Implementation hints

Buffering issues, particularly on CAN, have to be carefully considered. All nodes have to be able to keep up with PCER messages that can arrive about every millisecond (125000 bps / 120 bits per message)

- Don't do long linear searches to match an incoming Event ID to internal structures.
- Instead, consider just decoding the Event ID and setting a “process as available” bit, and then doing further work outside the CAN receive loop.

455 There is no requirement that events be processed in any particular order. Multiple PCER messages with the same Event ID may result in multiple actions inside a consumer, or may not, depending on the details of the consumer. There's a whole range of ways that events could be handled. A consumer might always handle them in the order received, and with the inter-event time intervals as received, down to the microsecond. Or it might do them in-order on a time-available basis, and generally be able to do them in 50 milliseconds. Or it might do them in the order that lets it get done fastest. Or whatever. Similar discussion goes for repeated events: Does the hardware have to do something twice?

460 That depends. In some cases, it's not even clear what it means. Specifying this is the node
 manufacturers job, because they have uses in mind, and they make nodes to meet those uses.
 OpenLCB just arranges to deliver events for ~~them~~ their processing. Node manufacturers should
 certainly specify the performance of their nodes. More knowledge is good, but the standard doesn't
 compel anything because it's hard to write a testable requirement.

465 The Standard doesn't make any statements about delivery performance because it's specific to the link-
 types, arrangement of OpenLCB nodes & links, etc that make up a specific OpenLCB network. In
 general, OpenLCB is going to provide sub-millisecond, millisecond, or perhaps even 10's of
 millisecond timing quality, for example. Nor can it guarantee absolute order of events. If Node A and
 Node B emit events at the same time on different segments, there's no guarantee of when node C will
 470 receive them, nor in what order, nor that Node D will receive it at the same time & in the same order as
 Node C.

3.4 Well known events

There are a small number of cases where a globally-allocated and reserved Event ID will simplify
 operation. These “well-known Event ID numbers” can be used to e.g. advertise that a node can provide
 a specific capability, or to tell locomotive control hardware to stop all trains instantly, etc.

475 For these to be useful, they not only have to be unique (so there are no collisions that accidentally
 trigger reactions to them), but they must also be well-known. All of these are assigned with the
 01.00.00.00.00.00 or 01.01.00.00.00.00 reserved IDs in their upper six bytes to ensure uniqueness and
 simplify recognition. The OpenLCB group maintains a central list of the individual assigned Event
 IDs.⁹

480 The Event IDs that start 01.00.00.00.00.00 are the automatically-routed events, a subset of the well-
 known events.

Except for the automatically-routed events, nodes producing or consuming well-known events must
 mention them when identifying the event IDs they produce and consume. Gateways may filter on these,
 but are not required to. For ease of implementation, a gateway may just pass all events with the
 485 common values of the top 6 bytes.

⁹See <http://www.openlcb.org/trunk/specs/EventIdAllocations.pdf> for specific definitions

Table of Contents

1	Introduction.....	1
2	Annotations to the Standard.....	2
2.1	Introduction.....	2
2.2	Intended Use.....	2
2.3	References and Context.....	2
2.4	Message Formats.....	2
2.4.1	Producer/Consumer Event Report (PCER).....	4
2.4.2	Identify Consumer.....	4
2.4.3	Consumer Identified.....	4
2.4.4	Consumer Range Identified.....	4
2.4.5	Identify Producer.....	4
2.4.6	Producer Identified.....	5
2.4.7	Producer Range Identified.....	5
2.4.8	Identify Events.....	5
2.4.9	Learn Event.....	6
2.5	States.....	6
2.6	Interactions.....	6
2.6.1	Event Transfer.....	7
2.6.2	Event Inquiry.....	7
2.6.3	Producer Inquiry.....	7
2.6.4	Consumer Inquiry.....	8
2.6.5	Teach/Learn Configuration.....	8
2.6.5.1	Blue/Gold Buttons Implementation Example.....	9
2.6.6	Resetting to Default.....	10
2.7	CAN Adaptation.....	10
2.7.1	CAN States.....	11
2.7.2	CAN Interactions.....	11
3	Background information.....	12
3.1	Gateways.....	12
3.2	Examples.....	12
3.2.1	Case 1: All nodes initialized before link established, event X is sent.....	13
3.2.2	Case 2: Node B2 comes up, after event X has been sent in the example above.....	13
3.2.3	Case 3: Node A1 comes up, then event X is sent.....	13
3.2.4	Case 4: Node A1 comes up, then event Y is sent.....	14
3.3	Node Implementation hints.....	14
3.4	Well known events.....	15